



2EL1590 – Cloud computing and distributed computing

Instructors: Francesca Bugiotti, Gianluca Quercini

Department: DÉPARTEMENT INFORMATIQUE

Language of instruction: FRANCAIS

Campus: CAMPUS DE PARIS - SACLAY

Workload (HEE): 60

On-site hours (HPE): 35,00

Elective Category : Fundamental Sciences

Advanced level : Yes

Description

Nowadays, the marketing strategies of most companies is based on the analysis of massive and heterogeneous data that needs a considerable amount of computational power. Instead of purchasing new hardware and software infrastructures, companies often resort to the computational and storage power offered by *cloud computing* platforms over the Internet.

The objective of this course is to present the fundamental principles of *distributed systems* and *distributed computing* that are at the heart of *cloud computing*.

The course will cover the principles of virtualization and containerization and the methods and tools used for distributed processing (for instance, *MapReduce*, *HDFS*, and *Spark*).

The course will also introduce advanced techniques and algorithms for the analysis of massive and heterogeneous data (PageRank, supervised learning, and *clustering*) and a brief introduction to some optimized Spark-compliant data formats (i.e., Parquet).

Quarter number

SG8

Prerequisites (in terms of CS courses)

Python programming, databases, basics of networking will be appreciated.

Syllabus

Introduction

- Cloud computing: motivation and terminology.
- Introduction to the public cloud providers (Amazon AWS, Microsoft Azure).
- Setup of a virtual machine on Microsoft Azure.



Virtualisation

- Virtualisation basics.
- Containerisation basics.
- Docker architecture.
- Images, containers, volumes and networks in Docker.
- Application deployment with Docker.

Multi-service applications and orchestration.

- Microservices architecture.
- Orchestration principles.
- Presentation of Kubernetes.
- Application deployment with Kubernetes.
- Application deployment in the cloud.

Cloud programming and software environments.

- Parallel computing, programming paradigms.
- Hadoop MapReduce.
- Apache Spark.
- Apache Parquet.

Data analysis.

- Cloud environments and data storage.
- Data distribution.
- Dataframes.

Class components (lecture, labs, etc.)

Introduction.

- Lecture : 3h

Virtualisation and containerisation.

- Lecture: 3h
- Tutorial : 3h

Multi-service applications.

- Lecture : 3h
- Tutorial : 3h



- Lab assignment (graded) : 3h

Cloud programming and software environments.

- Lecture : 9h
- Tutorial : 3h
- Lab assignment (graded) : 3h
- Exam: 2h

18h lecture, 9h tutorials, 6h lab assignments, 2h exam.

Grading

Written examination at the end of the course (MCQ + exercises) on the Evalmee platform (paperless exam).

- 2 lab assignments are graded.

Course support, bibliography

- Hwang, Kai, Jack Dongarra, and Geoffrey C. Fox. *Distributed and cloud computing: from parallel processing to the internet of things*. Morgan Kaufmann, 2013.
- Erl, T., Puttini, R., & Mahmood, Z. (2013). *Cloud computing: concepts, technology & architecture*. Pearson Education.
- Tel, G. (2000). *Introduction to distributed algorithms*. Cambridge university press.
- Miner, D., & Shook, A. (2012). *MapReduce Design Patterns: Building Effective Algorithms and Analytics for Hadoop and Other Systems*. O'Reilly Media, Inc.
- Karau, H., Konwinski, A., Wendell, P., & Zaharia, M. (2015). *Learning spark: lightning-fast big data analysis*. O'Reilly Media, Inc.
- Schenker, Gabriel. *Learn Docker - Fundamentals of Docker 18.x*. Packt Publishing,. Print.

Resources

Teaching staff: Francesca Bugiotti, Gianluca Quercini, Idir Ait Sadoune, Marc-Antoine Weisser, Arpad Rimmel

Maximum lab enrollment: 25 students

Software, number of licenses required: Use of free software



Learning outcomes covered on the course

At the end of this course, the students must be able to:

- Understand the fundamental concepts of cloud computing.
- Master the notion of virtualization and containerisation in the cloud.
- Be acquainted with the different cloud platforms.
- Use the distributed computing paradigms, such as MapReduce and Spark.
- Design distributed algorithms on data.

Description of the skills acquired at the end of the course

Operate all types of data, structured or unstructured, including big data.

- Conceive, design, implement and authenticate complex software.